

Programming Assignment V

Assigned: Oct. 28**Due:** Dec 9, at 11:59 PM (Checkpoint due Nov 18, at 11:59 PM)

1 Introduction

In this assignment, you will implement a code generator for Cool. When successfully completed, you will have a fully functional Cool compiler!

The code generator makes use of the AST constructed in PA3 and static analysis performed in PA4. Your code generator should produce MIPS assembly code that faithfully implements *any* correct Cool program. There is no error recovery in code generation—all erroneous Cool programs have been detected by the front-end phases of the compiler.

As with the static analysis assignment, this assignment has considerable room for design decisions. Your program is correct if the code it generates works correctly; how you achieve that goal is up to you. We will suggest certain conventions that we believe will make your life easier, but you do not have to take our advice. As always, explain and justify your design decisions in the README file. This assignment is about twice the amount of the code of the previous programming assignment, though they share much of the same infrastructure. **Start early!**

Critical to getting a correct code generator is a thorough understanding of both the expected behavior of Cool constructs and the interface between the runtime system and the generated code. The expected behavior of Cool programs are defined by the operational semantics for Cool given in Section 13 of the *Cool Reference Manual* (docs/cool-manual.pdf from the zip archive). Recall that this is only a specification of the meaning of the language constructs—not how to implement them. The interface between the runtime system and the generated code is given in *The Cool Runtime System* (docs/cool-runtime.pdf from the zip archive) See that document for a detailed discussion of the requirements of the runtime system on the generated code. There is a lot of information in this handout and the aforementioned documents, and you need to know most of it to write a correct code generator. *Please read thoroughly.*

You may work in a group of one or two people. The submit program will ask you to specify group members when you turn in your assignment.

2 Files and Directories

We provide all the necessary files in a downloadable zip available at <https://drive.google.com/open?id=0B0n20EI1U0EZTlgxMTR5S1JDSzg>. Please download and unzip wherever you plan to develop. All the scripts have been adapted so that any machine with Java 1.7, Apache Ant 1.9.6, and a Bash-compatible shell can run them. To install Apache Ant, follow the directions at <https://ant.apache.org/manual/index.html>.

You should not modify the following files (they are not symlinks): ASTConstants.java, ASTLexer.java, ASTParser.java, AbstractSymbol.java, AbstractTable.java, Cgen.java, ClassTable.java, Flags.java, IdSymbol.java, IdTable.java, IntTable.java, ListNode.java, build.xml, StringTable.java, SymbolTable.java, SymtabExample.java, TokenConstants.java, TreeNode.java, and Utilities.java. Almost all of these files have been described in previous assignments.

This is a list of the files (and folder) that you may want to modify. Most of the other files, you are probably familiar with from the previous assignments. See the README file for details about the additional files.

- **CgenClassTable.java** and **CgenNode.java**

These files provide an implementation of the inheritance graph for the code generator. You will need to complete **CgenClassTable** in order to build your code generator. You can use the provided code or replace it with your own inheritance graph from PA4. These skeletons are much larger than the ones for previous assignments. The skeletons provides three components of the code generator:

- functions to build the inheritance graph; (we supply this in case you didn't get this working for PA4)
- functions to emit global data and constants;

You should work to understand this code, and it will help you write the rest of the code generator.

- **StringSymbol.java**, **IntSymbol.java**, and **BoolConst.java**

These files provide support for Cool constants. You will need to complete the method for generating constant definitions.

- **cool-tree.java**

This file contains the definitions for the AST nodes. You will need to add code generation routines (**code(PrintStream)**) for Cool expressions in this file. The code generator is invoked by calling method **cgen(PrintStream)** of class **program**. You may add new methods, but do not modify the existing declarations.

- **TreeConstants.java**

As before, this file defines some useful symbol constants. Feel free to add your own as you see fit.

- **CgenSupport.java**

This file contains general support code for the code generator. You will find a number of handy functions here including ones for emitting MIPS instructions. Add to the file as you see fit, but don't change anything that's already there.

- **example.cl**

This file should contain a test program of your own design. Test as many features of the code generator as you can. You should write a correct Cool program which tests as many aspects of the code generator as possible. It should pass your code generator, and running MARS on the generated output should run the program correctly.

- **README**

This file will contain the write-up for your assignment. It is critical that you explain design decisions, how your code is structured, and why you believe your design is a good one (i.e., why it leads to a correct and robust program). It is part of the assignment to explain things in text as well as to comment your code.

- **tests**

is a directory containing seven test cases.

3 Instructions

Before beginning, ensure that Apache Ant is installed (via the directions at <https://ant.apache.org/manual/index.html>) and on your \$PATH.

To compile your compiler and code generator, cd into your PA5 directory and type:

```
"ant cgen"
```

To test your compiler, type:

```
"./mycoolc [-o output_filename.s] <file1.cl> <file2.cl> ..."
```

This command parses all the cool files given on the command line, passes them through the semantic checker, and then hands the program AST to your code generator.

To run your compiler on the file example.cl, type:

```
"ant test"
```

To run the produced code:

```
"./runmips file1.s" (or the output filename you chose)
```

To compare your compiler with the reference compiler, type:

```
"./compare-cgen <file1.cl>"
```

To turn in your work at checkpoint, copy the folder containing your PA5 code to the instr machine and type there:

```
"ant submit-clean"
```

```
"submit PA5-checkpoint"
```

To turn in your work finally, copy the folder containing your PA5 code to the instr machine and type there:

```
"ant submit-clean"
```

```
"submit PA5"
```

Be sure to submit all relevant source files.

In particular, you probably want to turn in cool-tree.java, TreeConstants.java, BoolConst.java, IntSymbol.java, StringSymbol.java, CgenNode.java, CgenClassTable.java, CgenSupport.java, example.cl, README.

You may turn in the assignment as many times as you like. However, only the last version will be retained for grading.

4 Design

Before continuing, we suggest you read *The Cool Runtime System* (docs/cool-runtime.pdf from the zip archive) to familiarize yourself with the requirements on your code generator imposed by the runtime system.

In considering your design, at a high-level, your code generator will need to perform the following tasks:

1. Determine and emit code for global constants, such as prototype objects.
2. Determine and emit code for global tables, such as the **class_nameTab**, the **class_objTab**, and the dispatch tables.
3. Determine and emit code for initialization method for each class.
4. Determine and emit code for each method definition.

There are many possible ways to write the code generator. One reasonable strategy is to perform code generation in two passes. The first pass decides the object layout for each class, particularly the offset at which each attribute is stored in an object. Using this information, the second pass recursively walks each feature and generates stack machine code for each expression.

There are a number of things you must keep in mind while designing your code generator:

- Your code generator must work correctly with the Cool runtime system, which is explained in the *Cool Runtime System* manual.
- You should have a clear picture of the runtime semantics of Cool programs. The semantics are described informally in the first part of the *Cool Reference Manual* (docs/cool-manual.pdf from the zip archive), and a precise description of how Cool programs should behave is given in Section 13 of the manual.
- You should understand the MIPS instruction set. An overview of MIPS operations is given in the MARS documentation, which is on the class web page.
- You should decide what invariants your generated code will observe and expect (i.e., what registers will be saved, which might be overwritten, etc). You may also find it useful to refer to information on code generation in the lecture notes and portions of the text, primarily ASU (The Dragon Book) Chapter 9.

You do *not* need to generate the same code as **coolc**. **Coolc** includes a very simple register allocator and other small changes that are not required for this assignment. The only requirement is to generate code that runs correctly with the runtime system.

4.1 Runtime Error Checking

The end of the Cool manual lists six errors that will terminate the program. Of these, your generated code should catch the first three—dispatch on void, case on void, and missing branch—and print a suitable error message before aborting. You may allow MARS to catch division by zero. Catching the last two errors—substring out of range and heap overflow—is the responsibility of the runtime system in **trap_handler.mars**. See Figure 4 of the *Cool Runtime System* manual for a listing of functions that display error messages for you.

4.2 Garbage Collection

To receive full credit for this assignment, your code generator must work correctly with the generational garbage collector in the Cool runtime system. The skeletons contain function **Cgen-ClassTable.codeSelectGc** that generate code that sets GC options from command line flags. The command-line flags that affect garbage collection are `-g`, `-t`, and `-T`. Garbage collection is disabled by default; the flag `-g` enables it. When enabled, the garbage collector not only reclaims memory, but also verifies that “-1” separates all objects in the heap, thus checking that the program (or the collector!) has not accidentally overwritten the end of an object. The `-t` and `-T` flags are used for additional testing. With `-t` the collector performs collections very frequently (on every allocation). The garbage collector does not directly use `-T`; in `coolc` the `-T` option causes extra code to be generated that performs more runtime validity checks. You are free to use (or not use) `-T` for whatever you wish.

For your implementation, the simplest way to start is not to use the collector at all (this is the default). When you decide to use the collector, be sure to carefully review the garbage collection interface described in the *Cool Runtime System* manual. Ensuring that your code generator correctly works with the garbage collector in *all* circumstances is not trivial.

5 Testing and Debugging

You will need a working scanner, parser, and semantic analyzer to test your code generator. You may use either your own components or the components from `reference-coolc`. By default, the `reference-coolc` components are used. To change that, replace the `lexer`, `parser`, and/or `semant` executable with your own scanner/parser. Even if you use your own components, it is wise to test your code generator with the `reference-coolc` scanner, parser, and semantic analyzer at least once because we will grade your project using `reference-coolc`’s version of the other phases.

You will run your code generator using `mycoolc`, a shell script that “glues” together the generator with the rest of compiler phases. Note that `mycoolc` takes a `-c` flag for debugging the code generator; using this flag merely causes `cgen_debug` (a static field of class `Flags`) to be set. Adding the actual code to produce useful debugging information is up to you. See the project `README` for details.

5.1 MARS

The executables `MARS` are simulators for MIPS architecture on which you can run your generated code. `MARS` has a GUI that lets you run MIPS assembly programs - it can also run them on the command line via the `runmips` script. The GUI has many features that allow you to examine the virtual machine’s state, including the memory locations, registers, data segment, and code segment of the program. You can also set breakpoints and single step your program. The documentation for `MARS` is in the course reader or on the course web page.

6 Checkpoint

This is a large project. Midway through the project we have a checkpoint, where we evaluate your progress. At that time, we want your code generation to work correctly on a very small set of programs. In particular, we want your code to work on programs with only one class -

Main, which inherits from IO - with one method `main` which immediately calls `out.int` with result an expression. This expression will not contain any `let`, `new`, or `case` expressions. An example program is in `tests/checkpoint.cl`. Once your compiler handles these test cases, you should submit the implementation by running “`submit PA5-checkpoint`”. The checkpoint submission is worth 40% of your grade for the project and will be evaluated on a variety of test cases.

7 Final Submission

Use “`submit PA5`” to submit your final result. Make sure to complete the following items before submitting to avoid any penalties.

- ☐ Include your write-up in `README`.
- ☐ Include your test cases that test your code generator in `example.cl`.
- ☐ Make sure all your code for the code generator is in
 - `cool-tree.java`, `CgenClassTable.java`, `CgenNode.java`, `CgenSupport.java`, `BoolConst.java`, `IntSymbol.java`, `StringSymbol.java`, `TreeConstants.java`, and additional `.java` files you may have added.
- ☐ Be sure to answer ‘yes’ to the submission prompt for files that contain your code.

8 Grading (out of 50 points)

The point breakdown for the PA5 final submission is as follows:

- 38 - for main autograder tests (20 points of which will be assigned during the checkpoint)
- 4 points for the `README`
 - 4 - thorough discussion of design decisions and choice of test cases; a few paragraphs of coherent English sentences should be fine
 - 2 - vague or hard to understand; omits important details
 - 0 - little to no effort
- 4 points - for `good.cl` and `bad.cl`
 - 4 - wide range of test cases added
 - 2 - added some tests, but the scope not sufficiently broad
 - 0 - little to no effort
- 4 points - for code cleanliness
 - 4 - code is mostly clean and well-commented
 - 2 - code is sloppy and/or poorly commented in places
 - 0 - little to no effort to organize and document code